

System Calls - Syntax

fork()

- The fork system call is used for creating a new process in Linux, and Unix systems, which is called the ***child process***, which runs concurrently with the process that makes the fork() call (parent process).
- After a new child process is created, both processes will execute the next instruction following the fork() system call.
- The child process uses the same pc(program counter), same CPU registers, and same open files which are used in the parent process.
- It takes no parameters and returns an integer value.

- Below are different values returned by fork().
- ***Negative Value***: The creation of a child process was unsuccessful.
- ***Zero***: Returned to the newly created child process.
- ***Positive value***: Returned to parent or caller. The value contains the process ID of the newly created child process.

```
#include <stdio.h>
#include <sys/types.h>;
#include <unistd.h>;
int main()
{
    // make two process which run same
    // program after this instruction
    fork();
    printf("Hello world!\n");
    return 0;
}
```

Output
Hello world
Hello world

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    fork();
    fork();
    fork();
    printf("Hello\n");
    return 0;
}
```

Output:-

Hello
Hello
Hello
Hello
Hello
Hello
Hello

The number of times 'hello' is printed is equal to the number of processes created. Total Number of Processes = 2^n , where n is the number of fork system calls. So here n = 3, $2^3 = 8$



exec() system call

- The exec family of functions replaces the current running process with a new process.
- It can be used to run a C program by using another C program.
- It comes under the header file **unistd.h**.
- There are many members in the exec family

- **execvp** : Using this command, the created child process does not have to run the same program as the parent process does.
- The **exec** type system calls allow a process to run any program files, which include a binary executable or a shell script .

```
int execvp (const *file, char *const argv[]);
```

- **file**: points to the file name associated with the file being executed.
- **argv**: is a null terminated array of character pointers.

We will have two .C files , **program1.c** and **program2.c**

And we will replace the **program1.c** with **program2.c** by calling execvp() function in program1.c .



```
//program2.c
```

```
#include<stdio.h>
#include<unistd.h>
int main()
{ int i;
    printf("I am program2.c called by execvp() ");
    printf("\n");
    return 0;
}
```

Compile the program and create an executable file program2 using
gcc program2.c -o program2

//program1.c

```
#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>

int main()
{ //A null terminated array of character pointers
    char *args[]={"./program2",NULL};
    execvp(args[0],args);
/*All statements are ignored after execvp() call as this whole
process(program1.c) is replaced by another process (program2.c)
*/
printf("Ending-----");
return 0;
}
```



Now create an executable file using gcc as

```
gcc program1.c -o program1
```

Now run it as

```
./program1
```

Output of the program

I am program2.c called by execvp()

getpid() and getppid()

```
// C Code to demonstrate getpid() and getppid()
```

```
#include<stdio.h>
```

```
#include <unistd.h>
```

```
int main()
```

```
{int pid;
```

```
pid = fork();
```

```
if (pid == 0)
```

```
{  printf( "\n Process id : %d " , getpid() );
```

```
    printf(" \n Process id : %d " , getppid());
```

```
}
```

```
return 0;
```


```
}
```



```
#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>
#include<sys/types.h>
int main(int argc, char **argv)
{ pid_t pid;
  pid = fork();
  if(pid==0)
  { printf("It is the child process and pid is %d\n",getpid());
    exit(0); }
  else if(pid > 0)
  { printf("It is the parent process and pid is %d\n",getpid()); }
  else
  { printf("Error while forking\n");
    exit(EXIT_FAILURE); }
  return 0; }
```

wait() system call

- A call to wait() blocks the calling process until one of its child processes exits or a signal is received.
- After child process terminates, parent ***continues*** its execution after wait system call instruction.



Syntax of wait()

```
// take one argument status  
// returns the process ID of dead child.  
pid_t wait(int *stat_loc);
```

wait takes argument NULL. There are also other means of arguments to wait, which is out of scope here.

```
// C program to demonstrate working of wait()
```

```
#include<stdio.h>
```

```
#include<stdlib.h>
```

```
#include<sys/wait.h>
```

```
#include<unistd.h>
```

```
int main()
```

```
{    pid_t cpid;
```

```
    if (fork() == 0)
```

```
        exit(0); /* terminate child */
```

```
    else
```

```
        cpid = wait(NULL); /* waiting parent */
```

```
        printf("Parent pid = %d\n", getpid());
```

```
        printf("Child pid = %d\n", cpid);
```

```
    return 0;
```

```
}
```



open() system call

Syntax

```
int open (path, flag)
```

Example:-

```
int fd1;
```

```
fd1= open("test.txt", O_RDONLY)
```

O_RDONLY stands for Read Only mode.

Return value

Returns the file descriptor , -1 upon failure

Close() system call

Closes a file

Syntax

```
int close(fp)
```

It takes the file pointer as input, return 0 on success, returns -1 if error.

// C program to illustrate open , close system Call

```
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    int fd1 = open("test.txt", O_RDONLY);
    if (fd1 < 0) {
        printf("Error\n");
        exit(1);
    }
    printf("opened the fd = %d\n", fd1);
    // Using close system Call
    if (close(fd1) < 0) {
        printf("Error\n");
        exit(1);
    }
    printf("closed the fd.\n");
}
```

stat() system call

- Returns information regarding a file

```
struct stat
{
    dev_t st_dev;          /* ID of device containing file */
    ino_t st_ino;          /* inode number */
    mode_t st_mode;        /* protection */
    nlink_t st_nlink;      /* number of hard links */
    uid_t st_uid;          /* user ID of owner */
    gid_t st_gid;          /* group ID of owner */
    dev_t st_rdev;         /* device ID (if special file) */
    off_t st_size;         /* total size, in bytes */
    blksize_t st_blksize;   /* blocksize for filesystem I/O */
    blkcnt_t st_blocks;    /* number of blocks allocated */
    time_t st_atime;       /* time of last access */
    time_t st_mtime;       /* time of last modification */
    time_t st_ctime;       /* time of last status change */};
```

```
#include<sys/stat.h>
#include<time.h>
#include<stdint.h>
int main()
{
    struct stat sfile;
    stat("even.c", &sfile);

    printf("Inode Number = %ld\n", sfile.st_ino);
    printf("Size = %ld\n", .....
```

The dirent structure is defined as:

```
#include<sys/dirent.h>
struct dirent
{ long d_ino; /* inode number */
  off_t d_off; /* offset to the next dirent */
  unsigned short d_reclen; /* length of this record */
  unsigned char d_type; /* type of file */
  char d_name[256]; /* filename*/  };
```

d_ino inode number of a file

d_offset offset of that directory entry in the actual file system directory

d_name name of the file or directory

d_reclen record length of this entry

d_type type of file(text file, block file, directory etc)

readdir()

- readdir() is a system call used to read into a directory.
- It returns a pointer to the dirent structure.

```
/* Unix program to read all regular files from a directory */
/printfile.c*/
#include<stdio.h>
#include<dirent.h>
#include<stdlib.h>
void main()
{ DIR *dirp;
struct dirent *dp;
if((dirp=opendir("/home/kumar"))==NULL) /* trying to open kumar
directory*/
{printf("\n cannot open");
exit(1); }
```

```
for(dp=readdir(dirp); dp!=NULL; dp=readdir(dirp))
{
    if(dp->d_type==DT_REG) /* printing only if file is a regular
file */
        printf("%s\n",dp->d_name);
}
closedir(dirp);
}
```

thanks

