

DEADLOCKS

DEADLOCKS

- Two processes each want to record a scanned document on a CD.
- Process *A* *requests permission to use the scanner and is granted it.*
- *Process B is programmed in such a way that it requests the CD recorder first and is also granted it.*
- Now *A asks for the CD recorder, but the request is denied until B releases it.*
- Instead of releasing the CD recorder *B asks for the scanner.*
- *At this point both processes are blocked and will remain so forever. This situation is called a **deadlock**.*

DEADLOCKS

- Deadlocks can occur on hardware resources or on software resources.
- Eg: In a database system, a program may have to lock several records it is using, to avoid race conditions. If process *A locks record R1 and process B locks record R2, and then each process tries to lock the other one's record, then there is a deadlock.*

DEADLOCKS

RESOURCES

- A resource is anything that can be used by only a single process at any instant of time.
- A resource can be a hardware device (e.g., a tape drive) or a piece of information (e.g., a locked record in a database).
- Some resources, several identical instances may be available, such as three tape drives.
- When several copies of a resource are available, any one of them can be used to satisfy any request for the resource.

DEADLOCKS

Preemptable and Nonpreemptable Resources

- A preemptable resource is one that can be taken away from the process owning it with no ill effects.
- Eg: Memory
- A non preemptable resource, is one that cannot be taken away from its current owner without causing the computation to fail.
- Eg: CD Recorder
- If a process has begun to burn a CD-ROM, suddenly taking the CD recorder away from it and giving it to another process will result in a garbled CD, CD recorders are not preemptable at an arbitrary moment.

DEADLOCKS

- Deadlocks involve nonpreemptable resources.
- Potential deadlocks that involve preemptable resources can usually be resolved by reallocating resources from one process to another.
- The sequence of events required to use a resource
 1. Request the resource.
 2. Use the resource.
 3. Release the resource.

DEADLOCKS

DEADLOCKS

- *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*
- A process requests resources; if the resources are not available at that time, the process enters a waiting state. Sometimes, a waiting process is never again able to change state, because the resources it has requested are held by other waiting processes. This situation is called a **deadlock**.

DEADLOCKS

Conditions for Deadlock

- A deadlock situation can arise if the following four conditions hold simultaneously in a system
 - 1. Mutual exclusion condition.**
 - 2. Hold and wait condition.**
 - 3. No preemption condition..**
 - 4.Circular wait condition.**

DEADLOCKS

- 1. Mutual exclusion:** At least one resource must be held in a nonsharable mode; that is, only one process at a time can use the resource. If another process requests that resource, the requesting process must be delayed until the resource has been released.
- 2. Hold and wait:** A process must be holding at least one resource and waiting to acquire additional resources that are currently being held by other processes.

DEADLOCKS

3. **No preemption:** Resources cannot be preempted; that is, a resource can be released only voluntarily by the process holding it, after that process has completed its task.
 4. **Circular wait.:** A set $\{P_0, P_1, \dots, P_n\}$ of waiting processes must exist such that P_0 is waiting for a resource held by P_1 , P_1 is waiting for a resource held by P_2 , ..., P_{n-1} is waiting for a resource held by P_n , and P_n is waiting for a resource held by P_0 .
- All four of these conditions must be present for a deadlock to occur. If one of them is absent, no deadlock is possible.

Resource-Allocation Graph

- Deadlocks can be described precisely in terms of a directed graph called a **system resource-allocation graph**.
- This graph consists of a set of vertices V and a set of edges E .
- *The set of vertices V is partitioned into two different types of nodes:*
 - $P = \{P1, P2, \dots, Pn\}$, *the set consisting of all the active processes in the system.*
 - $R = \{R1, R2, \dots, Rm\}$, *the set consisting of all resource types in the system.*

Resource-Allocation Graph

Request edge

- A directed edge $P_i \rightarrow R_j$ is called a **request edge**.
- It signifies that process P_i has requested an instance of resource type R_j and is currently waiting for that resource.


Assignment edge

- A directed edge $R_j \rightarrow P_i$ is called an **assignment edge**.
- It signifies that an instance of resource type R_j has been allocated to process P_i .

Resource-Allocation Graph

- Pictorially, we can represent each process P_i as a circle and each resource type R_j as a rectangle.
- Since resource type R_j may have more than one instance, we represent each instance as a dot within the rectangle.
- When process P_i requests an instance of resource type R_j , a request edge is inserted in the resource-allocation graph.
- When this request can be fulfilled, the request edge is transformed to an assignment edge.
- When the process no longer needs access to the resource, it releases the resource. As a result, the assignment edge is deleted.

Resource-Allocation Graph-Example



Resource-allocation graph.

Resource-Allocation Graph-Example

- The sets P , R , and E :

$$P = \{P1, P2, P3\}$$

$$R = \{R1, R2, R3, R4\}$$

$$E = \{P1 \rightarrow R1, P2 \rightarrow R3, R1 \rightarrow P2, R2 \rightarrow P2, R2 \rightarrow P1, R3 \rightarrow P3\}$$

- Resource instances:

- One instance of resource type $R1$
- Two instances of resource type $R2$
- One instance of resource type $R3$
- Three instances of resource type $R4$

- Process states:


- Process $P1$ is holding an instance of resource type $R2$ and is waiting for an instance of resource type $R1$.
- Process $P2$ is holding an instance of $R1$ and an instance of $R2$ and is waiting for an instance of $R3$.
- Process $P3$ is holding an instance of $R3$.

Resource-Allocation Graph


- If the graph contains no cycles, then no process in the system is deadlocked.
- If the graph does contain a cycle, then a deadlock may exist.
- If each resource type has exactly one instance, then a cycle implies that a deadlock has occurred.
- If the cycle involves only a set of resource types, each of which has only a single instance, then a deadlock has occurred. Each process involved in the cycle is deadlocked. In this case, a cycle in the graph is both a necessary and a sufficient condition for the existence of deadlock.
- If each resource type has several instances, then a cycle does not necessarily imply that a deadlock has occurred. In this case, a cycle in the graph is a necessary but not a sufficient condition for the existence of deadlock.

Resource-Allocation Graph

- If a resource-allocation graph does not have a cycle, then the system is *not in a deadlocked state*. *If there is a cycle, then the system may or* may not be in a deadlocked state.



Resource-allocation graph with a deadlock.



Resource-allocation graph with a cycle but no deadlock.